

Risk Checklists

- ▶ We're developing a system that auditors can use to monitor known risks that have been found in AI systems like those they're investigating.
- ▶ We have...
 - ▶ Persona and User Story
 - ▶ Mockups
 - ▶ Proposed API
- ▶ We'll be sharing these with our design partners (right), including Fairly.ai, who would like to integrate our data with their AI risk management platform.



Mockups

Discover

Submit

Welcome to the AIID

Discover Incidents

Risk Checklists

Table View

Entities

Taxonomies

Word Counts

Submit Incident Reports

Submission Leaderboard

Blog

Subscribe

AI INCIDENT DATABASE

English

Subscribe

Risk Checklist for "FinderBot"

Saving

Delete

Subscribe

Export

Agnes

Public Access

About System

Preview Markdown

Goals

Methods

Other Tags

Question Answering

Language Modeling

Transformer

Harm distribution basissace

Risks

Add Risk

Risks are surfaced automatically based on the tags applied to the system. They can also be added manually. Each risks is associated with a query for precedent incidents, which can be modified to suit your needs. You can subscribe both to new risks and new precedent incidents.

▼ Distributional Bias

Precedents Query

Distributional Bias

Question Answering

Language Modeling

Harm distribution

Precedents

Subscribe

Research Prototype AI, Delphi, Reportedly Gave Racially Biased Answers on Ethics

A publicly accessible research model that was trained via Reddit threads showed racially biased advice on moral dilemmas, allegedly demonstrating limitations of language-based models trained on moral judgments.

Question Answering

Language Modeling

Distributional Bias

Facebook translates 'good morning' into 'attack them', leading to arrest

Facebook's automatic language translation software incorrectly translated an Arabic post saying 'Good morning' into Hebrew saying

Risk Status

Severity

▲ Not mitigated

Risk Notes

Preview Markdown

▼ Gaming Vulnerability

NEW

▼ Detrimental Content

Unassessed

API

Our API design ergonomically exposes the risk monitoring features for use in third-party applications.

```
TAGS='GMF:Known AI Technology:Language Modeling__GMF:Known AI Goal:Question Answering'
```

```
curl -G "https://incidentdatabase.ai/api/riskManagement/v1/risks" \
```

```
--data-urlencode "tags=$TAGS&risks_added_since=2022-11-01"
```

```
{
  "risks": [
    {
      "tag": "GMF:Known AI Technical Failure:Misinformation Generation Hazard",
      "precedents": [
        {
          "incident_id": 470,
          "url": "https://incidentdatabase.ai/cite/470",
          "title": "Bing Chat Response Cited ChatGPT Disinformation Example",
          "description": "Reporters from TechCrunch[...]substantiate the disinformation.",
          "search_tags": [
            "GMF:Known AI Technology:Language Modeling",
            "GMF:Known AI Goal:Question Answering"
          ],
          "risk_tags": [
            "GMF:Known AI Technical Failure:Misinformation Generation Hazard"
          ]
        }
      ]
    }
  ]
}
```

Comments welcome!

Discover

Submit

Welcome to the AIID

Discover Incidents

Risk Checklists

Table View

Entities

Taxonomies

Word Counts

Submit Incident Reports

Submission Leaderboard

Blog

Subscribe

AI INCIDENT DATABASE

English

Subscribe

Risk Checklist for "FinderBot"

Saving

Delete

Subscribe

Export

Agnes

Public Access

About System

Preview Markdown

Goals

Methods

Other Tags

Question Answering

Language Modeling

Transformer

Harm distribution basissrace

Risks

Add Risk

Risks are surfaced automatically based on the tags applied to the system. They can also be added manually. Each risk is associated with a query for precedent incidents, which can be modified to suit your needs. You can subscribe both to new risks and new precedent incidents.

▼ Distributional Bias

Risk Status

Severity

Precedents Query

Precedents

Risk Notes

Distributional Bias

Question Answering

Language Modeling

Harm distribution

▼ Not mitigated

Subscribe

Preview Markdown

Research Prototype AI, Delphi, Reportedly Gave Racially Biased Answers on Ethics

A publicly accessible research model that was trained via Reddit threads showed racially biased advice on moral dilemmas, allegedly demonstrating limitations of language-based models trained on moral judgments.

Question Answering

Language Modeling

Distributional Bias

Facebook translates 'good morning' into 'attack them', leading to arrest

Facebook's automatic language translation software incorrectly translated an Arabic post saying "Good morning" into Hebrew saying

▼ Gaming Vulnerability

NEW

Unassessed

▼ Detrimental Content

NEW

Unassessed